

# 108-NLP @ NCTU

e-Lab 3

Boaz



# Good morning!

Please make sure your microphone/video is OFF

We'll take a break or two 

# Outline

Lab 2

N-Grams

Demo

Assignment

# Python is not JavaScript :)

Take advantage of Python!

<https://docs.python-guide.org/writing/style/>

# Lab 2 Part 1 - Proper Nouns

POS tag list:

CC	coordinating conjunction
CD	cardinal digit
DT	determiner
EX	existential there (like: "there is" ... think of it like "there exists")
FW	foreign word
IN	preposition/subordinating conjunction
JJ	adjective 'big'
JJR	adjective, comparative 'bigger'
JJS	adjective, superlative 'biggest'
LS	list marker 1)
MD	modal could, will
NN	noun, singular 'desk'
NNS	noun plural 'desks'
NNP	proper noun, singular 'Harrison'
NNPS	proper noun, plural 'Americans'
PDT	predeterminer 'all the kids'
POS	possessive ending parent\'s
PRP	personal pronoun I, he, she
PRP\$	possessive pronoun my, his, hers
RB	adverb very, silently,
RBR	adverb, comparative better
RBS	adverb, superlative best
RP	particle give up
TO	to go 'to' the store.
UH	interjection errrrrrrrm
VB	verb, base form take
VBD	verb, past tense took
VBG	verb, gerund/present participle taking
VBN	verb, past participle taken
VBP	verb, sing. present, non-3d take
VBZ	verb, 3rd person sing. present takes
WDT	wh-determiner which
WP	wh-pronoun who, what
WP\$	possessive wh-pronoun whose
WRB	wh-abverb where, when

MD	modal	could, will
NN	noun, singular	'desk'
NNS	noun plural	'desks'
NNP	proper noun, singular	'Harrison'
NNPS	proper noun, plural	'Americans'
PDT	predeterminer	'all the kids'
POS	possessive ending	parent\'s
PRP	personal pronoun	I, he, she

# Ask yourself: do the results MAKE SENSE?

[('San', 'Qualcomm'), 5354), (('U.', 'S.'), 5354), (('S.', 'Supreme'), 5354),  
(('Supreme', 'Court'), 5354), (('Lexmark', 'International'), 5354)]

[('U.', 'S.'), 1354280435), (('Donald', 'Trump'), 244930323), (('New', 'York'),  
193536141), (('Islamic', 'State'), 162978047), (('President', 'Donald'), 146460263)]

# Lab 2 Part 2

**Some news stories had missing content.**

Here are the titles:

- All The Action At The Golden Globes After Parties
- All The Looks At The MTV Movie & TV Awards Red Carpet
- 44 Of The Most Iconic Pictures Of President Barack Obama

**Why?**

One way to solve this problem:

```
news = pd.read_csv("https://bit.ly/nlp-buzzfeed").fillna('')
```

# This week's topic: 2-gram models

sentence = <s>, w1, w2, w3, ..., wi, wn, </s>

$p(\text{sentence}) \sim p(w_1 | \text{<s>}) p(w_2 | w_1) \dots p(\text{</s>} | w_n)$



# Google Books Ngram Viewer

<https://books.google.com/ngrams>

# Back to n-gram: Smoothing

Phenomena: Need exponentially large training data as  $n$  grows

- ◆ Infrequent terms as  $\langle \text{UNK} \rangle$
- ◆ Laplace's law:  $P(w_1, \dots, w_n) = (C(w_1, \dots, w_n) + 1) / (N + B)$ 
  - ◆ Referred as adding one, giving a little bit of the prob space to unseen events, for large  $N$  gives too much prob to the unseen  $n$ -gram
- ◆ Hold-out:  $P(w_1, \dots, w_n) = Tr / (NrN)$

# Resource

<https://web.stanford.edu/~jurafsky/slp3/3.pdf>

## 3.4.1 Laplace Smoothing

Laplace  
smoothing

The simplest way to do smoothing is to add one to all the bigram counts, before we normalize them into probabilities. All the counts that used to be zero will now have a count of 1, the counts of 1 will be 2, and so on. This algorithm is called **Laplace smoothing**. Laplace smoothing does not perform well enough to be used in modern n-gram models, but it usefully introduces many of the concepts that we see in other smoothing algorithms, gives a useful baseline, and is also a practical smoothing algorithm for other tasks like **text classification** (Chapter 4).

Let's start with the application of Laplace smoothing to unigram probabilities. Recall that the unsmoothed maximum likelihood estimate of the unigram probability of the word  $w_i$  is its count  $c_i$  normalized by the total number of word tokens  $N$ :

$$P(w_i) = \frac{c_i}{N}$$

add-one

Laplace smoothing merely adds one to each count (hence its alternate name **add-one smoothing**). Since there are  $V$  words in the vocabulary and each one was incremented, we also need to adjust the denominator to take into account the extra  $V$  observations. (What happens to our  $P$  values if we don't increase the denominator?)

$$P_{\text{Laplace}}(w_i) = \frac{c_i + 1}{N + V} \quad (3.20)$$

# Cross-Entropy and Perplexity

$$H(W) = -\frac{1}{N} \log P(w_1 w_2 \dots w_N) \quad (3.51)$$

**perplexity** The **perplexity** of a model  $P$  on a sequence of words  $W$  is now formally defined as the exp of this cross-entropy:

$$\begin{aligned} \text{Perplexity}(W) &= 2^{H(W)} \\ &= P(w_1 w_2 \dots w_N)^{-\frac{1}{N}} \\ &= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}} \\ &= \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}} \end{aligned} \quad (3.52)$$

# Demo

Full

<https://gist.github.com/bshmueli/42b80ef45fd1e704e8e93a76fb0a8210>

Short

<https://gist.github.com/bshmueli/b85471ad4f92b45ab91bab2b0f895c5a>

# Assignment

- Two parts (1 and 2).
  - This week we will use the following data:
    - [bit.ly/nlp-tweet-train](https://bit.ly/nlp-tweet-train) for training the model
    - [bit.ly/nlp-tweet-test](https://bit.ly/nlp-tweet-test) for evaluating the model
  - The data is in **JSON Lines** format. Each line is a JSON string.
  - You can read more about JSON Lines in [jsonlines.org](https://jsonlines.org)
- 
- TIP: you can use Pandas `read_json("filename", lines=True)` to read the data

# Part 1

Build a 2-gram model for the Twitter train data:

- For tokenization, first convert to lowercase, then use the NLTK TweetTokenizer. That's it!
- Don't forget to add the '<s>' and '</s>' tokens to each tweet.
- Vocabulary: only frequent terms (terms appearing **3 or more** times in the train data). Replace infrequent terms (appearing only 1 or 2 times in the train data) by '<UNK>' in both train and test data
- Use Laplace smoothing (identical to the one we used in the class)

## Output (2 numbers)

- Average perplexity for:  
Train data tweets  
Test data tweets

# Part 2

Build a **bi-directional** 2-gram model by training on the Twitter train data:

- Build 2-gram forward model (**identical to part 1**)
- Build 2-gram backward model (**identical to part 1**, but in the opposite direction)
- **NEW! See next slide for more information about the bi-directional language model**
- $\gamma$ : a parameter between 0 and 1

## Output (3 numbers)

- Print the  $\gamma$  that minimizes the perplexity of the Twitter test data (0.05 resolution)
- Average perplexity (at the optimal  $\gamma$ ) for:
  - Training data tweets
  - Testing data tweets



# Bi-directional Language Model

For the regular forward-direction bi-gram model we use:

$$p_f(w_1, \dots, w_n) \approx p(w_2|w_1) \cdot \dots \cdot p(w_i|w_{i-1}) \dots \cdot p(w_n|w_{n-1})$$

where

$$p(w_i|w_{i-1}) = \frac{C(w_{i-1}, w_i) + 1}{C(w_{i-1}) + |V|}$$

Similarly, for the backward bi-gram model:

$$p_b(w_1, \dots, w_n) \approx p(w_1|w_2) \cdot \dots \cdot p(w_{n-1}|w_n)$$

For the bidirectional model, you will estimate the probability of each word using a linear combination of the probabilities of the forward and backward bi-grams for each word:

$$p(w_i|w_{i-1}, w_{i+1}) = \gamma \cdot p(w_i|w_{i-1}) + (1 - \gamma) \cdot p(w_i|w_{i+1})$$

The probability for the whole tweet will be:

$$p_{bd}(w_1, \dots, w_n) = p(w_2|w_1, w_3) \cdot \dots \cdot p(w_i|w_{i-1}, w_{i+1}) \cdot p(w_{n-1}|w_{n-2}, w_n)$$

Note that  $w_1$  is  $\langle s \rangle$  and  $w_n$  is  $\langle /s \rangle$ .

# SOP for assignment submission

- ☐ Open your Python notebook
- ☐ Rename your notebook to lab3-<studentID>.ipynb
- ☐ “Restart and run all”
- ☐ Make sure the output is correct
- ☐ Download the Python file (File, Download .py)
- ☐ **Share the Google colab using “Get shareable link”; Copy link**
- ☐ Submit both link and file in e3:
  - ☐ Paste link into e3 (link starts with **<https://colab.research.google.com/>**...)
  - ☐ Upload file into e3 (file name is lab3-<studentID>.py)

# Resources

<https://web.stanford.edu/~jurafsky/slp3/3.pdf>

[How to find the perplexity of a corpus](#)

[https://rstudio-pubs-static.s3.amazonaws.com/115676\\_ab6bb49748c742b88127e8b5ce3e1298.html](https://rstudio-pubs-static.s3.amazonaws.com/115676_ab6bb49748c742b88127e8b5ce3e1298.html)

<https://www.youtube.com/watch?v=CTYqkWU8cBc>

<http://www.cs.columbia.edu/~mccollins/courses/nlp2011/notes/lm.pdf>

<https://courses.cs.washington.edu/courses/csep517/18au/slides/csep517au18-LanguageModels.pdf>

<https://stackoverflow.com/questions/54941966/how-can-i-calculate-perplexity-using-nltk>